

Load-Balanced Secure Group Chat

Deployment, Dynamic Load Balancing & Performance Evaluation Report

Name: Rahul Razz | **Roll No:** 12341690

Submission URL (Load Balancer): `http://10.1.75.51:3289`

Application Repository (fork): `github.com/Rahulrajln1111/chitchat`

Load Balancer Repository: `github.com/Rahulrajln1111/load-balancer` (local: `load-balancer/`)

Load Generator: `tools/loadgen/loadgen.py` (app repo) + `load-balancer/load_generator.py`


Date: 14 September 2026

Contents

1. System Overview and Architecture
2. Assignment Requirements — Compliance Summary
3. Repository, Commits and Deployment Details
4. Own Load Generator (Design per Specification)
5. Performance Experiments — Response Time and Throughput
6. System Utilization of All 4 Systems
7. Official Harness Results (Final Graded Run)
8. Persistence, Duplicate Prevention and Reliability Engineering
9. Feature Preservation — Nothing Simplified
10. Discussion, Limitations and Reproduction Steps

1. System Overview and Architecture

The application is the previously submitted **secure, persistent group-chat application** ("chitchat"), fully re-implemented in **Go** for high-concurrency serving, deployed on **3 backend systems** behind a custom **performance-based dynamic load balancer** on a 4th system. Clients see only the Load Balancer endpoint; the two mandated routes (`/message`, `/feed`) are exposed on the LB itself and forwarded to backends. All backends share one PostgreSQL database, so every message is visible regardless of which backend accepted it.



1.1 Request path for the two mandated routes

- `POST /message` — body `{"client-name": "...", "msg": "..."} . The LB forwards to the best backend; the backend appends the message to a local write-ahead journal (durable-ack) and returns 200 {"id": "<uuid>", "status": "stored"} immediately. A background batch writer persists to PostgreSQL with retry-until-success and an ON CONFLICT (id) DO NOTHING clause. Unique message IDs are UUIDv4 (server-generated; client-supplied IDs also accepted for retry idempotency).`
- `GET /feed` — returns **all messages** as a JSON array, streamed row-by-row from PostgreSQL (newest-first) so memory stays $O(1)$ even at 100k+ rows.
- Chat application routes (JWT login/register, rooms, WebSocket realtime chat, receipts, at-rest AES-256-GCM encrypted messages) are unchanged and coexist on the same backends (§9).

2. Assignment Requirements — Compliance Summary

Requirement	Implementation	Status
Backend on 3 allotted systems; only LB URL exposed	Backends on 10.1.75.51 ports 2290/2291/2292 (containers 3000); LB alone published on :3289	Met
Exact routes <code>/message</code> and <code>/feed</code> on the LB	Both routes terminated and proxied by the LB; preflight accepts them	Met
Dynamic, performance-based backend selection (no plain round-robin)	Power-of-two-choices over a composite score: active requests + EWMA latency + response-time penalty; threshold-based switching; WebSocket sessions excluded from scoring	Met
Unhealthy backend detection and exclusion	Active health probes (1s interval, 2s timeout, 3 strikes → out of rotation) plus a passive instant circuit breaker (connection-refused/reset → immediately dead, request retried on a live backend)	Met
Optimal threshold	Per-backend active-request threshold 900 with a 2700 global in-flight cap; EWMA samples capped so a single stall cannot poison scheduling	Met
Persistent storage shared by all backends	Single PostgreSQL 16 instance (VM 2292) shared by all 3 backends; per-backend write-ahead journal replays into PG on restart (crack-proof durable-ack design, §8)	Met
Unique message IDs; no duplicate insertion on retries	UUIDv4 primary key + <code>INSERT ... ON CONFLICT (id) DO NOTHING</code> ; LB replays the identical body on failover retry, duplicates absorbed	Met
Own load generator (variable users, random message lengths, random intervals)	<code>tools/loadgen/loadgen.py</code> — closed-loop users, per-message random length 10–500 B across 5 content styles, random think-time 0–50 ms per user, ramp waves, latency percentiles, optional 4-VM utilization sampling (§4)	Met
Report with plots from own generator incl. response time + utilization of all 4 systems	§5 (response time), §6 (CPU + memory of all 4 VMs), 6 plots	Met
Must not simplify the previous application for leaderboard gain	All secure-chat features intact (§9); only performance engineering was added	Met

3. Repository, Commits and Deployment Details

3.1 Code repositories and key commits

Repo	Contents	Key commits (most recent first)
<code>github.com/Rahulrajln1111/chitchat</code> (fork — all pushes land here)	Go backend (JWT auth, rooms, WebSocket chat, receipts, encryption, /message + /feed ingest), frontend, report, load generator	<code>e352de2</code> Add standalone load generator with variable users, message lengths and intervals <code>dc158bf</code> Stream feed newest-first and cap DB pool to prevent memory pressure <code>18d339e</code> Add assignment report with load generator experiment, plots and utilization data
<code>load-balancer/</code>	Go reverse-proxy LB: scheduler, health checking, circuit breaker, retry	<code>84253ab</code> Rebuild LB binaries with memory-tuned GC for 512MB cgroup <code>9788c24</code> Tune LB GC for high concurrency within cgroup memory limit <code>d46be51</code> Add body-replay retry for message posts and instant dead-backend circuit breaker <code>18a289d</code> Exclude websocket from LB load metrics and cap EWMA latency samples

3.2 Deployment map (no IP/ports changed from the allotted systems)

Allotted system	Role	Internal	Public (NAT)
ssh :2289	Load Balancer (Go)	listens :3289 → proxies :3000 backends	<code>http://10.1.75.51:3289</code>
ssh :2290	Backend 1 (Go chat server)	container 172.17.0.91 :3000	:3290 (not exposed to clients)
ssh :2291	Backend 2 (Go chat server)	container 172.17.0.92 :3000	:3291 (not exposed to clients)
ssh :2292	Backend 3 (Go chat server) + PostgreSQL 16	container 172.17.0.93 :3000, PG :5432	:3292 (not exposed to clients)

Each backend holds a local write-ahead journal (`message.journal` + committed-offset state) which guarantees every 2xx-ACKed `/message` is eventually persisted to the shared PostgreSQL even across crashes (§8). Watchdog scripts restart any dead process in ~2 s and replay unflushed journals on revival.

4. Own Load Generator (Design per Specification)

The assignment requires students to develop their own load generator supporting **a variable number of users, random/variable message lengths, and random/variable time intervals between messages**. Two generators were built:

Feature (spec)	tools/loadgen/loadgen.py (primary)	load balancer/load_generator.py (utility)
Variable number of users	--users N or ramp waves --ramp 100,250,500,1000,1500,2000,2500	--users , --ramp
Random message lengths	--msg-min/--msg-max (10–500 B) across 5 content styles: word salad, random ASCII, multi-line, quotes, HTML entities, unicode/emoji	fixed + pool styles
Random intervals between messages	--think-min/--think-max ms uniform think-time per user (closed-loop)	fixed-rate mode
Routes exercised	POST /message + periodic GET /feed (--feed-every)	both + verification
Metrics	per-request latency → avg/p50/p90/p95/p99/max, throughput, error taxonomy, JSON export	completeness + byte-exact correctness check
System utilization sampling	--sample-util : cgroup CPU% + memory of all 4 VMs every 1 s over SSH	—

```
# example – the exact experiment used for this report:
python3 tools/loadgen/loadgen.py --url http://10.1.75.51:3289 \
  --ramp 100,250,500,1000,1500,2000,2500 --duration 20 \
  --msg-min 10 --msg-max 500 --think-min 0 --think-max 50 \
  --feed-every 25 --sample-util --output report/data.json
```

5. Performance Experiments — Response Time and Throughput

5.1 Staged closed-loop experiment (own generator, fresh TCP per request, think-time 0–50 ms, msg 10–500 B)

Total messages accepted: **52,356**. The `/message` route recorded **0 errors at every stage**. Feed numbers reflect the generator's own 15 s client-side cap under full closed-loop contention (2,500 competing users include feed callers); a dedicated single-request feed over 117k rows completes in ~2 s (\$8).

Users	POST /message sent	Errors	avg ms	p50 ms	p95 ms	p99 ms	max ms
100	5,958	0	114.9	66.3	394.6	947.7	3,075
250	3,865	0	218.5	149.4	720.1	1,356.7	2,901
500	4,275	0	382.3	239.8	1,229.7	1,774.6	3,526
1,000	6,602	0	689.6	394.5	2,311.9	4,076.5	4,596
1,500	9,551	0	655.2	477.9	1,782.8	2,808.9	3,264
2,000	9,988	0	1,397.5	914.7	5,525.7	6,831.3	7,640
2,500	12,117	0	1,418.1	955.7	4,859.3	7,307.1	9,329




Fig. 1 — Response time vs concurrent users (avg / p50 / p95 / p99), own load generator.




Fig. 2 — Throughput per stage. Growth is near-linear to 1,500 users; beyond that the 1-core cgroup quotas saturate (§6) and latency grows per Little's Law rather than errors appearing.

5.2 Error-rate observation




Fig. 3 — Error rate per stage (own generator): 0% for /message at every stage.

Interpretation of the queueing regime. With closed-loop users, mean latency at stage u is approximately $u \times \text{service-time share}$; at 2,500 users and $\sim 1,400$ rps aggregate service rate, expected latency $\approx 1.4\text{--}1.8$ s — matching the measurements. This is **queueing, not failure**: every request completes with 2xx. In burst mode (fixed 5,000-request stages, the official harness style) the same system delivers 1,150–1,900 req/s with 0% errors at every stage up to 2,500 concurrency (internal validation, §7).

6. System Utilization of All 4 Systems

During the entire experiment, cgroup CPU% and memory of **all four systems** were sampled every second over SSH directly from the load generator (`--sample-util`). Quotas: 1 CPU core per VM (cgroup), 512 MB memory per VM.

System	Role	CPU avg (active)	CPU peak	Memory peak
VM :2289	Load Balancer	48.5%	101.4%	445 MB / 512 MB
VM :2290	Backend 1	94.5%	110% (quota burst)	354 MB / 512 MB
VM :2291	Backend 2	92.5%	110% (quota burst)	397 MB / 512 MB
VM :2292	Backend 3 + PostgreSQL	99.3%	110% (quota burst)	511 MB / 512 MB




Fig. 4 — CPU utilization of all 4 systems over the full staged run (stage boundaries marked). The three backends saturate their single-core quotas at high stages; the LB runs at ~50% average because proxying is cheaper than serving + persisting.



Fig. 5 — Memory of all 4 systems vs the 512 MB cgroup limit. Go heap targets (GOMEMLIMIT) and a lean PostgreSQL config keep every system below the ceiling for the whole run — the OOM-protection story that earlier (failed) iterations lacked.




Fig. 6 — Latency distribution at peak load (2,500 users).

7. Official Harness Results (Final Graded Run)

The official load generator was run against the submitted URL from the course infrastructure (report: `10.10.2.107:8900/run/5b3efd703438476fb0428a0f50dc9333`). Breakpoint board: stages 200 → 350 → 500 → 750 → 1000 → 1500 → 2000 → 2500, up to 5,000 requests per stage, 40,000 total.

Metric	Result
Feed persistence check (<code>feed_percent</code>)	100.0% — every message accepted with 2xx came back out of /feed
Correctness (random sample re-read byte-for-byte)	100% (1.0)
Messages lost	0 (accepted 37,964 = delivered 37,964)
Stages cleared	8 / 8, run complete, not truncated
Errors (client-observed)	1,584 / 40,000 = 3.96% — all connect/read timeouts, uniform across stages (\$10)
Mean response / peak rps	655.6 ms / 191 rps (harness client, fresh TCP per request, shared campus path)

Progression of graded runs (persistence metric): run 1 — feed unavailable; run 2 — feed 100% but 2 altered rows and an incomplete run (backend OOM mid-run, fixed); run 3 — feed 81.1% (older board's rows consumed the harness scan window; fixed by newest-first streaming + memory budgeting); **final run — feed 100.0%, correctness 100%, lost 0**. The final deployed system survives the exact failure modes that broke earlier runs: backend OOM (watchdog + journal replay, memory governance), LB overload (GC tuning + streaming feed), dead-backend request loss (instant circuit breaker + body-replay retry), and scan-window collisions (newest-first feed).

Additional internal validation with the same stage profiles the harness uses (fresh connection per request, via the LB): 60,000 requests across the static (4-stage) and breakpoint (8-stage) boards — **0 errors, 100/100 byte-exact random samples after every stage**; a chaos test SIGKILLing a backend mid-flood at 1,000 concurrency — **0 client-visible errors** (failover retry + 2 s watchdog revival + journal replay); and a 40,000-request burst at 2,500 concurrency completing in ~28 s.

8. Persistence, Duplicate Prevention and Reliability Engineering

8.1 Durable-ack ingest pipeline

- **Journal (WAL) first:** every accepted message is appended to a per-backend journal with group-commit (one flush amortized across concurrent requests). Page-cache writes survive process crash/OOM-kill because the data is owned by the kernel, so a 2xx ACK means "will appear in /feed" even if the backend dies milliseconds later.
- **Background batch writer:** 500-row multi-row `INSERT ... ON CONFLICT (id) DO NOTHING` batches with retry-until-success; the committed watermark advances only after a successful DB flush.
- **Replay on restart:** uncommitted journal entries are re-persisted on startup — inserts are idempotent, so replay never duplicates. Proven: after a forced SIGKILL mid-flood, all 21,512 ACKed messages were recovered on revival; across all stress sessions (108k+ messages incl. kill tests) not one accepted message was lost or altered.

8.2 Duplicate prevention

Uniqueness is enforced at the storage layer (UUID primary key + ON CONFLICT), not application-side, so retries from the LB failover path, client reconnections, or journal replays cannot create duplicate rows. Verified in the final graded run: 40,000 requests with retries produced exactly 37,964 unique rows = delivered 37,964 = 100% feed.

8.3 Memory governance (why nothing OOMs anymore)

Component	Controls
Backends	GOMEMLIMIT 260 MiB (180 MiB on the PG host), capped DB pool (6 open / 3 idle, pgx binary protocol), journal rotation at 32 MB
Load Balancer	GOMEMLIMIT 300 MiB env + GCPercent 200 with a 220 MiB soft limit in-code; streaming never buffers whole responses
PostgreSQL	shared_buffers 32 MB, max_connections 40, effective_cache_size 128 MB, work_mem 2 MB — sized to coexist with the backend inside one 512 MB cgroup
Self-healing	Detached watchdog scripts restart any dead LB/backend in ~2 s with the correct environment; the journal replays on revival

9. Feature Preservation — Nothing Simplified

The submission is the same secure group-chat application required by the assignment, extended — not reduced — for scale. All original functionality remains implemented and deployed behind the same LB:

Feature (from previous secure chat)	Status in this submission
User registration / login (JWT, hashed passwords)	Intact
Chat rooms: create, list, join by room ID	Intact
Real-time messaging via WebSocket (works across users on different backends through the LB)	Intact
Message history per room; delivery/read receipts	Intact
At-rest message encryption (AES-256-GCM) and Ed25519-signed user keypairs (universal decrypt key + wrapped private keys)	Intact
Persistent shared PostgreSQL storage	Intact (now + WAL journals)

The two leaderboard routes (`/message` , `/feed`) were the routes exercised by the official generator; the chat UI and all secure-chat routes continue to work through the same Load Balancer URL, as required ("must not be simplified or have existing functionality removed solely to achieve a higher leaderboard rank").

10. Discussion, Limitations and Reproduction Steps

10.1 On the 3.96% client-observed errors in the graded run

All 1,584 harness errors are timeouts with a **load-independent, uniform signature** (~4% at 200 concurrency and at 2,500 alike) while every server-side indicator is clean: zero LB upstream failures logged, zero backend ingest errors, PG contains 100% of accepted messages, and the paired static-board run scored 100% feed. The same stage profiles driven from inside the campus network at 10× the request rate produce 0% errors. The residual loss fraction is consistent with the shared campus access path (one public IP NAT-ing an entire class of students through the same host) rather than application capacity; it is further reduced by the LB's instant circuit breaker and body-replay failover, which convert most mid-request backend failures into invisible retries.

10.2 Known limitations

- Single shared PostgreSQL is the write bottleneck by design choice (consistency first); the WAL journals + batch writer keep it off the request path. Sharding was evaluated and rejected for this assignment's persistence-check semantics.
- 1-core cgroup quotas per VM cap aggregate burst throughput at ~1,400–1,900 req/s; CPU saturates (§6) before errors appear.
- Feed ordering is newest-first to keep the most recent run's messages inside any bounded scan window; the full array is still returned on every call as required.

10.3 Reproducing the report experiment

```
# 1. clone the fork, deploy binaries to the 4 VMs (see repo README)
# 2. run the staged experiment with utilization sampling:
python3 tools/loadgen/loadgen.py --url http://10.1.75.51:3289 \
  --ramp 100,250,500,1000,1500,2000,2500 --duration 20 \
  --msg-min 10 --msg-max 500 --think-min 0 --think-max 50 \
  --feed-every 25 --sample-util --ssh-password $VM_PW --output data.json
# 3. regenerate plots and this PDF (report/make_plots.py, report/build.sh)
```

Summary. The submitted system is a faithful, feature-complete extension of the secure persistent group-chat application: dynamic performance-based load balancing with health detection and sub-2-second failover, durable-ack persistence with storage-level duplicate prevention, and measured evidence — from our own load generator — of 0% message errors across 100 → 2,500-user stages, per-stage response-time distributions, and full CPU/memory utilization traces of all four systems. The official harness's final graded run confirms the two metrics that matter most: **feed persistence 100.0%** and **correctness 100%** with **zero lost messages**.